

rl

November 16, 2019

1 IE 534 HW: Reinforcement Learning

v1, Designed by TIANQI WU, Fall 2019 at UIUC

In this assignment, we will experiment with the (deep) reinforcement learning algorithms covered in the lecture. In particular, you will implement variants of the popular DQN (Deep Q-Network) (1) and A2C (Advantage Actor-Critic) (2) algorithms (by the same first author! orz), and test your implementation on both a small example (CartPole problem) and an Atari game (Breakout game). We focus on model-free algorithms rather than model-based ones, because neural nets are easier applicable and more popular nowadays in the model-free setting. (When the system dynamic is known or can be easily inferred, model-based can sometimes do better.)

The assignment breaks into **three parts**:

- **In Part I** (50 pts), you basically need to follow the instructions in this notebook to do a little bit of coding. We'll be able to see if your code trains by testing against the CartPole environment provided by the OpenAI gym package. We'll generate some plots that are required for grading.
- **In Part II** (40 pts), you'll copy your code onto Blue Waters (or actually any good server..), and run a much larger-scale experiment with the Breakout game. Hopefully, you can teach the computer to play Breakout in less than half a day! Share your final game score in this notebook. **This part will take at least a day. Please start early!!**
- **In Part III** (10 pts), you'll be asked to think about a few questions. These questions are mostly open-ended. Please write down your thoughts on them.

Finally, after you finished everything in this notebook (**code snippets C1-C5, plots P1-P5, question answers Q1-Q5**), please save the notebook, and export to a PDF (or an HTML file), and submit:

1. the **.ipynb notebook and exported .pdf/.html file**, PDF is preferred (I usually do File -> Print Preview -> use Chrome's Save as PDF);
2. your code (**Algo.py, Model.py files**);
3. job artifacts (**.log files** only, pytorch models and images not required)

to Compass 2g for grading.

PS: Remember to save your notebook occasionally as you work through it!

References

- (1) Mnih, V., Kavukcuoglu, K., Silver, D., Rusu, A.A., Veness, J., Bellemare, M.G., Graves, A., Riedmiller, M., Fidjeland, A.K., Ostrovski, G. and Petersen, S., 2015. Human-level control through deep reinforcement learning. *Nature*, 518(7540), p.529.
- (2) Mnih, V., Badia, A.P., Mirza, M., Graves, A., Lillicrap, T., Harley, T., Silver, D. and Kavukcuoglu, K., 2016, June. Asynchronous methods for deep reinforcement learning. In *International conference on machine learning* (pp. 1928-1937).
- (3) A useful tutorial: <https://spinningup.openai.com/en/latest/>
- (4) Useful code references: <https://github.com/deepmind/bsuite>; <https://github.com/openai/baselines>; <https://github.com/astooke/rlpyt>;

First of all, **enter your NetID here** in the cell below: Your NetID: twu38

1.1 Part I: DQN and A2C on CartPole

This part is designed to run on your own local laptop/PC.

Before you start, there are some python dependencies: `pytorch`, `gym`, `numpy`, `multiprocessing`, `matplotlib`. Please install them correctly. You can install `pytorch` following instruction here <https://pytorch.org/get-started/locally/>. The code is compatible with PyTorch 0.4.x ~ 1.x. PyTorch 1.1 with cuda 10.0 worked for me (`conda install pytorch==1.1.0 torchvision==0.3.0 cudatoolkit=10.0 -c pytorch`).

Please **always** run the code cell below each time you open this notebook, to make sure `gym` is installed and to enable `autoreload` which **allows code changes to be effective immediately**. So if you changed `Algo.py` or `Model.py` but the test codes are not reflecting your changes, restart the notebook kernel and run this cell!!

```
In [1]: # install openai gym
        # pip install gym
        # enable autoreload
        %load_ext autoreload
        %autoreload 2
```

1.1.1 1.1 Code Structure

The code is structured in 5 python files:

- `Main.py`: contains the main entry point and training loop
- `Model.py`: constructs the torch neural network modules
- `Env.py`: contains the environment simulations interface, based on openai gym
- `Algo.py`: implements the DQN and A2C algorithms
- `Replay.py`: implements the experience replay buffer for DQN
- `Draw.py`: saves some game snapshots to jpeg files

Some parts of the code `Model.py` and `Algo.py` are left blank for you to complete. You are not required to modify the other parts (unless, of course, you want to boost the performance!). This is kind of a minimalist implementation, and might be different from the other code on the internet in details. You're welcomed to improve it, after you've finished all the required things of this assignment.

1.1.2 1.2 OpenAI gym and CartPole environment

OpenAI developed python package `gym` a while ago to facilitate RL research. `gym` provides a common interface between the program and the environments. For instance, the code cell below will create the `CartPole` environment. A window will show up when you run the code. The goal is to keep adjusting the cart so that the pole stays in its upright position.

A demo video from OpenAI:

`gym` also provides interface to Atari games. However, installing package `atari-py` is not easy on Windows/Mac, so we won't demonstrate it here. More info: <http://gym.openai.com/docs/>.

```
In [2]: import time
import gym
env = gym.make('CartPole-v1')
env.reset()
for _ in range(200):
    env.render()
    state, reward, done, _ = env.step(env.action_space.sample()) # take a random action
    if done: break
    time.sleep(0.15)
env.close()
```

1.1.3 1.3 Deep Q Learning

A little recap on DQN. We learned from lecture that Q-Learning is a model-free reinforcement learning algorithm. It falls into the off-policy type algorithm since it can utilize past experiences stored in a buffer. It also falls into the (approximate) dynamic programming type algorithm, since it tries to learn an optimal state-action value function using time difference (TD) errors. Q Learning is particularly interesting because it exploits the optimality structure in MDP. It's related to the Hamilton-Jacobi-Bellman equation in classical control.

For MDP

$$M = (S, A, P, r, \gamma)$$

where S is the state space, A is the action space, P is the transition dynamic, $r(s, a)$ is a reward function $S \times A \mapsto R$, and γ is the discount factor.

The tabular case (when S, A are finite), Q-Learning does the following value iteration update repeatedly when collecting experience (s_t, a_t, r_t) (η is the learning rate):

$$Q^{new}(s_t, a_t) \leftarrow Q^{old}(s_t, a_t) + \eta \left(r_t + \gamma \max_{a' \in A} Q^{old}(s_t, a') - Q^{old}(s_t, a_t) \right).$$

With function approximation, meaning model $Q(s, a)$ with a function $Q_\theta(s, a)$ parameterized by θ , we arrive at the Fitted Q Iteration (FQI) algorithm, or better known as Deep Q Learning if the function class is neural networks. Q-Learning with neural network as function approximator was

known long ago, but it was only recently (year 2013) that DeepMind made this algorithm actually work on Atari games. Deep Q Learning iteratively optimize the following objective:

$$\theta_{new} \leftarrow \arg \min_{\theta} \mathbb{E}_{(s,a,r,s') \sim D} \left(r + \gamma \max_{a' \in A} Q_{\theta_{old}}(s', a') - Q_{\theta}(s, a) \right)^2.$$

Therefore, with a batch of $\{(s^i, a^i, r^i, s^i)\}_{i=1}^N$ sampled from the replay buffer, we can build a loss function L in pytorch:

$$L(\theta) = \frac{1}{N} \sum_{i=1}^N \left(r^i + \gamma \max_{a' \in A} Q_{\theta_{old}}(s^i, a') - Q_{\theta}(s^i, a^i) \right)^2,$$

and run the usual gradient descent on θ with a pytorch optimizer.

Exploration Exploration, as the word suggests, refers to explore novel unvisited states in RL. The FQI (or DQN) needs an exploratory datasets to work well. The common way to produce exploratory dataset is through randomization, such as the ϵ -greedy exploration strategy we will implement in this assignment. - ϵ -greedy exploration:

At training iteration it , the agent chooses to play

$$a = \begin{cases} \arg \max_a Q_{\theta}(s, a) & \text{with probability } 1 - \epsilon_{it} , \\ \text{a random action } a \in A & \text{with probability } \epsilon_{it} . \end{cases}$$

And ϵ_{it} is annealed, for example, linearly from 1 to 0.01 as training progresses until iteration it_{decay} :

$$\epsilon_{it} = \max \left\{ 0.01, 1 + (0.01 - 1) \frac{it}{it_{decay}} \right\}.$$

Two Caveats

1. There's an improvement on DQN called Double-DQN with the following loss L , which has shown to be empirically more stable than the original DQN loss described above. You may want to implement the improved one in your code:

$$L(\theta) = \frac{1}{N} \sum_{i=1}^N \left(r^i + \gamma Q_{\theta_{old}}(s^i, \arg \max_{a' \in A} Q_{\theta}(s^i, a')) - Q_{\theta}(s^i, a^i) \right)^2.$$

2. Huber loss (a.k.a smooth L1 loss) is commonly used to reduce the effect of extreme values:

$$L(\theta) = \frac{1}{N} \sum_{i=1}^N \text{Huber} \left(r^i + \gamma Q_{\theta_{old}}(s^i, \arg \max_{a' \in A} Q_{\theta}(s^i, a')) - Q_{\theta}(s^i, a^i) \right)$$

You can look up the pytorch document here: <https://pytorch.org/docs/stable/nn.functional.html#smooth-l1-loss>

C1 (5 pts): Complete the code for the two layered fully connected network class `TwoLayerFCNet` in file `Model.py` And run the cell below to test the output shape of your module.

```
In [3]: ## Test code
        from Model import TwoLayerFCNet
        import torch
        net = TwoLayerFCNet(n_in=4, n_hidden=16, n_out=5)
        x = torch.randn(10, 4)
        y = net(x)
        assert y.shape == (10, 5), "ERROR: network output has the wrong shape!"
        print ("Output shape test passed!")
```

Output shape test passed!

C2 (5 pts): Complete the code for ϵ -greedy exploration strategy in function `DQN.act` in file `'Algo.py'` And run the cell below to test it.

```
In [4]: ## Test code
        from Algo import DQN
        class Nothing: pass
        dummy = Nothing()
        dummy.eps_decay = 500000

        dummy.num_act_steps = 0
        eps = DQN.compute_epsilon(dummy)
        assert abs( eps - 1.0 ) < 0.01, "ERROR: compute_epsilon at t=0 should be 1 but got %f!"

        dummy.num_act_steps = 250000
        eps = DQN.compute_epsilon(dummy)
        assert abs( eps - 0.505 ) < 0.01, "ERROR: compute_epsilon at t=250000 should around .505"

        dummy.num_act_steps = 500000
        eps = DQN.compute_epsilon(dummy)
        assert abs( eps - 0.01 ) < 0.01, "ERROR: compute_epsilon at t=500000 should be .01 but got %f!"

        dummy.num_act_steps = 600000
        eps = DQN.compute_epsilon(dummy)
        assert abs( eps - 0.01 ) < 0.01, "ERROR: compute_epsilon after t=500000 should be .01 but got %f!"
        print ("Epsilon-greedy test passed!")
```

Epsilon-greedy test passed!

C3 (10 pts): Complete the code for computing the loss function in `DQN.train` in file `Algo.py` And run the cell below to verify your code decreases the loss value in one iteration.

```
In [5]: import numpy as np
        from Algo import DQN
```

```

class Nothing: pass
dummy_obs_space, dummy_act_space = Nothing(), Nothing()
dummy_obs_space.shape = [10]
dummy_act_space.n = 3

dqn = DQN(dummy_obs_space, dummy_act_space, batch_size=2)

for t in range(3):
    dqn.observe([np.random.randn(10).astype('float32')], [np.random.randint(3)],
                [(np.random.randn(10).astype('float32'), np.random.rand(), False, None)

    b = dqn.replay.cur_batch
    loss1 = dqn.train()
    dqn.replay.cur_batch = b
    loss2 = dqn.train()

    print (loss1, '>', loss2, '?')
    assert loss2 < loss1, "DQN.train should reduce loss on the same batch"

    print ("DQN.train test passed!")

```

parameters to optimize: [('fc1.weight', torch.Size([128, 10]), True), ('fc1.bias', torch.Size([128, 10]), False)]

0.05873310565948486 > 0.056507933884859085 ?
DQN.train test passed!

P1 (10 pts): Run DQN on CartPole and plot the learning curve (i.e. averaged episodic reward against env steps). Your code should be able to achieve >150 averaged reward in 10000 iterations (20000 simulation steps) in only a few minutes. This is a good indication that the implementation is correct. It's ok that the curve is not always monotonically increasing because of randomness in training.

```

In [6]: %run Main.py \
        --niter 10000 \
        --env CartPole-v1 \
        --algo dqn \
        --nproc 2 \
        --lr 0.001 \
        --train_freq 1 \
        --train_start 100 \
        --replay_size 20000 \
        --batch_size 64 \
        --discount 0.996 \
        --target_update 1000 \
        --eps_decay 4000 \
        --print_freq 200 \
        --checkpoint_freq 20000 \

```

```

--save_dir cartpole_dqn \
--log log.txt \
--parallel_env 0

```

```

Namespace(algo='dqn', batch_size=64, checkpoint_freq=20000, discount=0.996, ent_coef=0.01, env=
observation space: Box(4,)
action space: Discrete(2)
running on device cpu
parameters to optimize: [('fc1.weight', torch.Size([128, 4]), True), ('fc1.bias', torch.Size([

```

```

obses on reset: 2 x (4,) float32

```

iter	200	loss	0.01	n_ep	15	ep_len	21.4	ep_rew	21.38	raw_ep_rew	21.38	env_step
iter	400	loss	0.00	n_ep	29	ep_len	27.0	ep_rew	27.01	raw_ep_rew	27.01	env_step
iter	600	loss	0.00	n_ep	44	ep_len	29.3	ep_rew	29.33	raw_ep_rew	29.33	env_step
iter	800	loss	0.00	n_ep	63	ep_len	21.6	ep_rew	21.60	raw_ep_rew	21.60	env_step
iter	1000	loss	0.00	n_ep	86	ep_len	16.0	ep_rew	15.95	raw_ep_rew	15.95	env_step
iter	1200	loss	0.03	n_ep	110	ep_len	17.6	ep_rew	17.64	raw_ep_rew	17.64	env_step
iter	1400	loss	0.05	n_ep	137	ep_len	14.7	ep_rew	14.67	raw_ep_rew	14.67	env_step
iter	1600	loss	0.05	n_ep	165	ep_len	15.3	ep_rew	15.29	raw_ep_rew	15.29	env_step
iter	1800	loss	0.02	n_ep	186	ep_len	16.4	ep_rew	16.38	raw_ep_rew	16.38	env_step
iter	2000	loss	0.04	n_ep	211	ep_len	15.5	ep_rew	15.48	raw_ep_rew	15.48	env_step
iter	2200	loss	0.06	n_ep	241	ep_len	13.9	ep_rew	13.92	raw_ep_rew	13.92	env_step
iter	2400	loss	0.05	n_ep	266	ep_len	16.6	ep_rew	16.58	raw_ep_rew	16.58	env_step
iter	2600	loss	0.06	n_ep	281	ep_len	26.5	ep_rew	26.55	raw_ep_rew	26.55	env_step
iter	2800	loss	0.04	n_ep	301	ep_len	20.6	ep_rew	20.57	raw_ep_rew	20.57	env_step
iter	3000	loss	0.06	n_ep	319	ep_len	24.1	ep_rew	24.07	raw_ep_rew	24.07	env_step
iter	3200	loss	0.08	n_ep	329	ep_len	29.9	ep_rew	29.91	raw_ep_rew	29.91	env_step
iter	3400	loss	0.01	n_ep	335	ep_len	45.9	ep_rew	45.93	raw_ep_rew	45.93	env_step
iter	3600	loss	0.14	n_ep	339	ep_len	67.6	ep_rew	67.63	raw_ep_rew	67.63	env_step
iter	3800	loss	0.08	n_ep	346	ep_len	67.1	ep_rew	67.13	raw_ep_rew	67.13	env_step
iter	4000	loss	0.02	n_ep	351	ep_len	59.1	ep_rew	59.12	raw_ep_rew	59.12	env_step
iter	4200	loss	0.22	n_ep	357	ep_len	61.8	ep_rew	61.77	raw_ep_rew	61.77	env_step
iter	4400	loss	0.08	n_ep	360	ep_len	76.3	ep_rew	76.28	raw_ep_rew	76.28	env_step
iter	4600	loss	0.06	n_ep	364	ep_len	100.8	ep_rew	100.85	raw_ep_rew	100.85	env_step
iter	4800	loss	0.07	n_ep	366	ep_len	108.9	ep_rew	108.87	raw_ep_rew	108.87	env_step
iter	5000	loss	0.07	n_ep	368	ep_len	113.0	ep_rew	113.02	raw_ep_rew	113.02	env_step
iter	5200	loss	0.03	n_ep	370	ep_len	133.4	ep_rew	133.39	raw_ep_rew	133.39	env_step
iter	5400	loss	0.19	n_ep	372	ep_len	149.3	ep_rew	149.27	raw_ep_rew	149.27	env_step
iter	5600	loss	0.07	n_ep	374	ep_len	160.2	ep_rew	160.25	raw_ep_rew	160.25	env_step
iter	5800	loss	0.07	n_ep	377	ep_len	158.4	ep_rew	158.44	raw_ep_rew	158.44	env_step
iter	6000	loss	0.02	n_ep	378	ep_len	162.5	ep_rew	162.50	raw_ep_rew	162.50	env_step
iter	6200	loss	0.14	n_ep	379	ep_len	179.8	ep_rew	179.85	raw_ep_rew	179.85	env_step
iter	6400	loss	0.05	n_ep	381	ep_len	183.7	ep_rew	183.67	raw_ep_rew	183.67	env_step
iter	6600	loss	0.22	n_ep	383	ep_len	204.8	ep_rew	204.81	raw_ep_rew	204.81	env_step
iter	6800	loss	0.02	n_ep	385	ep_len	199.2	ep_rew	199.18	raw_ep_rew	199.18	env_step
iter	7000	loss	0.09	n_ep	387	ep_len	199.6	ep_rew	199.55	raw_ep_rew	199.55	env_step
iter	7200	loss	0.03	n_ep	389	ep_len	206.9	ep_rew	206.95	raw_ep_rew	206.95	env_step
iter	7400	loss	0.06	n_ep	391	ep_len	199.9	ep_rew	199.87	raw_ep_rew	199.87	env_step

```

iter 7600 |loss 0.51 |n_ep 393 |ep_len 214.0 |ep_rew 214.02 |raw_ep_rew 214.02 |env_step
iter 7800 |loss 0.09 |n_ep 394 |ep_len 214.6 |ep_rew 214.62 |raw_ep_rew 214.62 |env_step
iter 8000 |loss 0.02 |n_ep 395 |ep_len 219.5 |ep_rew 219.46 |raw_ep_rew 219.46 |env_step
iter 8200 |loss 0.09 |n_ep 397 |ep_len 226.3 |ep_rew 226.30 |raw_ep_rew 226.30 |env_step
iter 8400 |loss 0.43 |n_ep 399 |ep_len 220.4 |ep_rew 220.38 |raw_ep_rew 220.38 |env_step
iter 8600 |loss 0.28 |n_ep 401 |ep_len 226.9 |ep_rew 226.87 |raw_ep_rew 226.87 |env_step
iter 8800 |loss 0.07 |n_ep 403 |ep_len 224.7 |ep_rew 224.74 |raw_ep_rew 224.74 |env_step
iter 9000 |loss 0.03 |n_ep 405 |ep_len 220.8 |ep_rew 220.81 |raw_ep_rew 220.81 |env_step
iter 9200 |loss 0.05 |n_ep 407 |ep_len 230.0 |ep_rew 230.04 |raw_ep_rew 230.04 |env_step
iter 9400 |loss 0.10 |n_ep 408 |ep_len 227.5 |ep_rew 227.54 |raw_ep_rew 227.54 |env_step
iter 9600 |loss 0.04 |n_ep 410 |ep_len 224.8 |ep_rew 224.76 |raw_ep_rew 224.76 |env_step
iter 9800 |loss 0.04 |n_ep 412 |ep_len 220.6 |ep_rew 220.61 |raw_ep_rew 220.61 |env_step
save checkpoint to cartpole_dqn/9999.pth

```

```
In [7]: import matplotlib.pyplot as plt
```

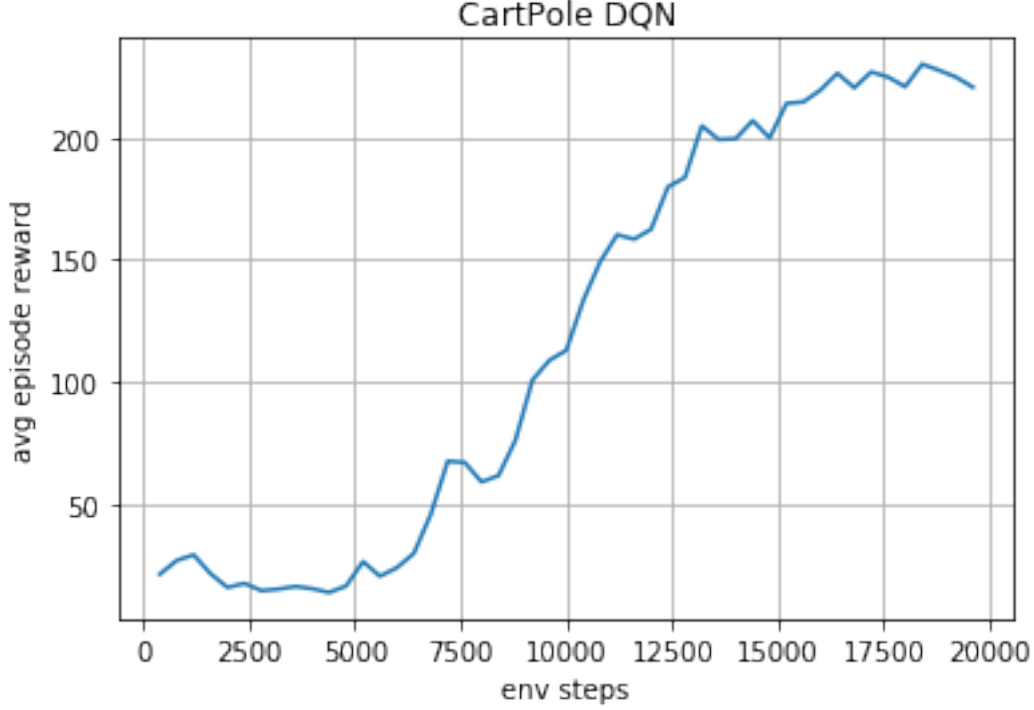
```

def plot_curve(logfile, title=None):
    lines = open(logfile, 'r').readlines()
    lines = [l.split() for l in lines if l[:4] == 'iter']
    steps = [int(l[13]) for l in lines]
    rewards = [float(l[11]) for l in lines]
    plt.plot(steps, rewards)
    plt.xlabel('env steps'); plt.ylabel('avg episode reward'); plt.grid(True)
    if title: plt.title(title)
    plt.show()

```

The log is saved to 'cartpole_dqn/log.txt'. Let's plot the running averaged episode reward curve during training:

```
In [8]: plot_curve('cartpole_dqn/log.txt', 'CartPole DQN')
```

1.1.4 1.4 Actor-Critic Algorithm

Policy gradient methods are another class of algorithms that originated from viewing the RL problem as a mathematical optimization problem. Recall that the objective of RL is to maximize the expected cumulative reward the agent gets, namely

$$\max_{\pi} J(\pi) := \mathbb{E}_{(s_t, a_t, r_t) \sim D^{\pi}} \left[\sum_{t=0}^{\infty} \gamma^t r_t \right]$$

where D^{π} is the distribution of trajectories induced by policy π , and inside the expectation is the random variable representing the discounted cumulative reward and J is the reward (or cost) functional. Essentially, we want to optimize the policy π .

The most straightforward way is to run gradient update on the parameter θ of a *parameterized* policy π_{θ} . One such algorithm is the so-called Advantage Actor-Critic (A2C). A2C is an on-policy policy optimization type algorithm. While collecting on-policy data, we iteratively run gradient ascent:

$$\theta_{new} \leftarrow \theta_{old} + \eta \hat{\nabla}_{\theta} J(\pi_{\theta_{old}})$$

with a Monte Carlo estimate $\hat{\nabla}_{\theta} J$ of the true gradient $\nabla_{\theta} J$. The true gradient writes as (by Policy Gradient Theorem and some manipulations):

$$\nabla_{\theta} J(\pi_{\theta_{old}}) = \mathbb{E}_{(s_t, a_t, r_t) \sim D^{\pi_{\theta_{old}}}} \sum_{t=0}^{\infty} \left(\nabla_{\theta} \log \pi_{\theta_{old}}(s_t, a_t) \left(\sum_{t'=t}^{\infty} \gamma^{t'-t} r_{t'} - V^{\pi_{\theta_{old}}}(s_t) \right) \right).$$

The quantity in the inner-most parentheses $A(s_t, a_t) = Q(s_t, a_t) - V(s_t) = (\mathbb{E} \sum_{t'=t}^{\infty} \gamma^{t'-t} r_{t'}) - V(s_t)$ is called the *Advantage* function (not very rigorously speaking...). That's why it's called **Advantage** Actor-Critic. More on A2C: <https://arxiv.org/abs/1506.02438>.

And the Monte Carlo estimate of the gradient is

$$\hat{\nabla}_{\theta} J(\pi_{\theta_{old}}) = \frac{1}{NT} \sum_{i=1}^N \sum_{t=0}^T \left(\nabla_{\theta} \log \pi_{\theta_{old}}(s_t^i, a_t^i) \left(\sum_{t'=t}^T \gamma^{t'-t} r_{t'}^i - V_{\phi_{old}}(s_t^i) \right) \right)$$

where $V_{\phi_{old}}$ is introduced as a *parameterized* estimate for $V^{\pi_{\theta_{old}}}$, which can also be a neural network. So V_{ϕ} is the **critic** and π_{θ} is the **actor**. We can construct a specific loss function in pytorch that gives $\hat{\nabla}_{\theta} J$. $V_{\phi_{old}}$ is trained with SGD on a L2 loss function. It's further common practice to add an entropy bonus loss term to encourage maximum entropy solution, to facilitate exploration and avoid getting stuck in local minima. We shall clarify these loss functions in the following summarization.

Summarizing a variant of the A2C algorithm:

For many iterations repeat: 1. Collect N independent trajectories $\{(s_t^i, a_t^i, r_t^i)_{t=0}^T\}_{i=1}^N$ by running policy π_{θ} for maximum T steps; 2. Compute the loss function for the policy parameter θ :

$$L_{policy}(\theta) = \frac{1}{NT} \sum_{i=1}^N \sum_{t=0}^T \left(\log \pi_{\theta}(s_t^i, a_t^i) \left(\sum_{t'=t}^T \gamma^{t'-t} r_{t'}^i - V_{\phi}(s_t^i) \right) \right)$$

Compute the entropy term for θ :

$$L_{entropy}(\theta) = \frac{1}{NT} \sum_{i=1}^N \sum_{t=0}^T \left(- \sum_{a \in A} \pi_{\theta}(s_t^i, a) \log \pi_{\theta}(s_t^i, a) \right)$$

Compute the loss for value function parameter ϕ :

$$L_{value}(\phi) = \frac{1}{NT} \sum_{i=1}^N \sum_{t=0}^T \left(\sum_{t'=t}^T \gamma^{t'-t} r_{t'}^i - V_{\phi}(s_t^i) \right)^2$$

3. Use pytorch auto-differentiation and optimizer to do one gradient step on (θ, ϕ) with the overall loss:

$$L(\theta, \phi) = -L_{policy} - \lambda_{ent} L_{entropy} + \lambda_{val} L_{value}$$

where λ_{ent} and λ_{val} are coefficients to balances the loss terms.

C4 (10 pts): Complete the code for computing the advantage, entropy and loss function in A2C.train in file Algo.py

In []:

P2 (10 pts): Run A2C on CartPole and plot the learning curve (i.e. averaged episodic reward against training iteration). Your code should be able to achieve >150 averaged reward in 10000 iterations (40000 simulation steps) in only a few minutes. This is a good indication that the implementation is correct.

```
In [9]: %run Main.py \
        --niter 10000 \
        --env CartPole-v1 \
        --algo a2c \
        --nproc 4 \
        --lr 0.001 \
        --train_freq 16 \
        --train_start 0 \
        --batch_size 64 \
        --discount 0.996 \
        --value_coef 0.01 \
        --print_freq 200 \
        --checkpoint_freq 20000 \
        --save_dir cartpole_a2c \
        --log log.txt \
        --parallel_env 0
```

```
Namespace(algo='a2c', batch_size=64, checkpoint_freq=20000, discount=0.996, ent_coef=0.01, env=
observation space: Box(4,)
action space: Discrete(2)
running on device cpu
shared net = False, parameters to optimize: [('fc1.weight', torch.Size([128, 4]), True), ('fc1
```


```
obses on reset: 4 x (4,) float32
```

iter	200	loss	0.87	n_ep	38	ep_len	18.4	ep_rew	18.43	raw_ep_rew	18.43	env_step
iter	400	loss	0.94	n_ep	77	ep_len	25.3	ep_rew	25.26	raw_ep_rew	25.26	env_step
iter	600	loss	0.79	n_ep	116	ep_len	20.7	ep_rew	20.73	raw_ep_rew	20.73	env_step
iter	800	loss	0.77	n_ep	155	ep_len	22.9	ep_rew	22.95	raw_ep_rew	22.95	env_step
iter	1000	loss	0.65	n_ep	186	ep_len	23.4	ep_rew	23.38	raw_ep_rew	23.38	env_step
iter	1200	loss	0.67	n_ep	219	ep_len	26.6	ep_rew	26.57	raw_ep_rew	26.57	env_step
iter	1400	loss	0.67	n_ep	246	ep_len	29.5	ep_rew	29.50	raw_ep_rew	29.50	env_step
iter	1600	loss	0.80	n_ep	268	ep_len	34.3	ep_rew	34.31	raw_ep_rew	34.31	env_step
iter	1800	loss	0.85	n_ep	288	ep_len	45.0	ep_rew	44.96	raw_ep_rew	44.96	env_step
iter	2000	loss	1.02	n_ep	303	ep_len	44.8	ep_rew	44.79	raw_ep_rew	44.79	env_step
iter	2200	loss	0.64	n_ep	319	ep_len	58.1	ep_rew	58.08	raw_ep_rew	58.08	env_step
iter	2400	loss	0.93	n_ep	332	ep_len	56.1	ep_rew	56.14	raw_ep_rew	56.14	env_step
iter	2600	loss	0.95	n_ep	346	ep_len	51.2	ep_rew	51.18	raw_ep_rew	51.18	env_step
iter	2800	loss	1.03	n_ep	359	ep_len	54.8	ep_rew	54.75	raw_ep_rew	54.75	env_step
iter	3000	loss	0.93	n_ep	372	ep_len	59.4	ep_rew	59.40	raw_ep_rew	59.40	env_step
iter	3200	loss	0.59	n_ep	387	ep_len	56.0	ep_rew	56.01	raw_ep_rew	56.01	env_step
iter	3400	loss	0.66	n_ep	396	ep_len	73.8	ep_rew	73.80	raw_ep_rew	73.80	env_step
iter	3600	loss	0.99	n_ep	403	ep_len	79.0	ep_rew	79.02	raw_ep_rew	79.02	env_step
iter	3800	loss	0.67	n_ep	415	ep_len	79.0	ep_rew	79.05	raw_ep_rew	79.05	env_step

iter	4000	loss	0.99	n_ep	423	ep_len	81.7	ep_rew	81.68	raw_ep_rew	81.68	env_step
iter	4200	loss	0.40	n_ep	431	ep_len	90.3	ep_rew	90.26	raw_ep_rew	90.26	env_step
iter	4400	loss	0.96	n_ep	434	ep_len	127.3	ep_rew	127.26	raw_ep_rew	127.26	env_step
iter	4600	loss	0.35	n_ep	440	ep_len	131.0	ep_rew	130.98	raw_ep_rew	130.98	env_step
iter	4800	loss	0.45	n_ep	448	ep_len	123.7	ep_rew	123.73	raw_ep_rew	123.73	env_step
iter	5000	loss	0.57	n_ep	458	ep_len	94.1	ep_rew	94.08	raw_ep_rew	94.08	env_step
iter	5200	loss	0.30	n_ep	464	ep_len	100.4	ep_rew	100.37	raw_ep_rew	100.37	env_step
iter	5400	loss	0.85	n_ep	471	ep_len	125.7	ep_rew	125.67	raw_ep_rew	125.67	env_step
iter	5600	loss	1.07	n_ep	475	ep_len	131.1	ep_rew	131.09	raw_ep_rew	131.09	env_step
iter	5800	loss	0.85	n_ep	480	ep_len	142.6	ep_rew	142.61	raw_ep_rew	142.61	env_step
iter	6000	loss	0.49	n_ep	491	ep_len	106.0	ep_rew	106.02	raw_ep_rew	106.02	env_step
iter	6200	loss	0.07	n_ep	496	ep_len	113.4	ep_rew	113.37	raw_ep_rew	113.37	env_step
iter	6400	loss	0.88	n_ep	504	ep_len	106.4	ep_rew	106.40	raw_ep_rew	106.40	env_step
iter	6600	loss	0.19	n_ep	508	ep_len	122.5	ep_rew	122.50	raw_ep_rew	122.50	env_step
iter	6800	loss	-0.07	n_ep	514	ep_len	148.4	ep_rew	148.37	raw_ep_rew	148.37	env_step
iter	7000	loss	0.73	n_ep	517	ep_len	164.9	ep_rew	164.85	raw_ep_rew	164.85	env_step
iter	7200	loss	0.90	n_ep	521	ep_len	191.0	ep_rew	191.01	raw_ep_rew	191.01	env_step
iter	7400	loss	0.03	n_ep	526	ep_len	190.3	ep_rew	190.25	raw_ep_rew	190.25	env_step
iter	7600	loss	-0.03	n_ep	527	ep_len	194.5	ep_rew	194.53	raw_ep_rew	194.53	env_step
iter	7800	loss	0.61	n_ep	533	ep_len	181.2	ep_rew	181.21	raw_ep_rew	181.21	env_step
iter	8000	loss	0.64	n_ep	538	ep_len	177.5	ep_rew	177.53	raw_ep_rew	177.53	env_step
iter	8200	loss	-0.07	n_ep	544	ep_len	156.2	ep_rew	156.22	raw_ep_rew	156.22	env_step
iter	8400	loss	0.74	n_ep	547	ep_len	162.9	ep_rew	162.94	raw_ep_rew	162.94	env_step
iter	8600	loss	0.13	n_ep	553	ep_len	180.4	ep_rew	180.35	raw_ep_rew	180.35	env_step
iter	8800	loss	0.75	n_ep	557	ep_len	175.9	ep_rew	175.87	raw_ep_rew	175.87	env_step
iter	9000	loss	0.14	n_ep	562	ep_len	163.8	ep_rew	163.84	raw_ep_rew	163.84	env_step
iter	9200	loss	-0.07	n_ep	568	ep_len	167.1	ep_rew	167.14	raw_ep_rew	167.14	env_step
iter	9400	loss	0.73	n_ep	570	ep_len	165.6	ep_rew	165.58	raw_ep_rew	165.58	env_step
iter	9600	loss	0.76	n_ep	573	ep_len	183.6	ep_rew	183.61	raw_ep_rew	183.61	env_step
iter	9800	loss	0.71	n_ep	576	ep_len	212.7	ep_rew	212.71	raw_ep_rew	212.71	env_step

save checkpoint to cartpole_a2c/9999.pth

In [10]: plot_curve('cartpole_a2c/log.txt', 'CartPole A2C')



Now let's play a little bit with the trained agent. The neural net parameters are saved to the `cartpole_dqn` and `cartpole_a2c` folders. The cell below will open a window showing one episode play.

```
In [11]: import time
import gym
import Algo
env = gym.make('CartPole-v1')
agent = Algo.ActorCritic(env.observation_space, env.action_space)
agent.load('cartpole_a2c/9999.pth')
state = env.reset()
for _ in range(120):
    env.render()
    state, reward, done, _ = env.step(agent.act([state])[0])
    if done: break
    time.sleep(0.1)
env.close()
```

shared net = False, parameters to optimize: [('fc1.weight', torch.Size([128, 4]), True), ('fc1

1.2 Part II: Solve the Atari Breakout game

In this part, you'll train your agent to play Breakout with the BlueWaters cluster. I have provided the job scripts for you. Please upload your `Algo.py` and `Model.py` completed in **Part I** to your BlueWaters folder. And submit the following two jobs respectively:

```
qsub run_dqn.pbs
qsub run_a2c.pbs
```

The jobs are set to run for at most **14 hours**. **Please start early!!** You might be able to reach the desired score (≥ 200 reward) before 14 hours - You can stop the training early if you wish. Then please collect the resulting `breakout_dqn/log.txt` and `breakout_a2c/log.txt` files into the same folder as this Jupyter notebook's. Rename them as `log_breakout_dqn.txt` and `log_breakout_a2c.txt`.

BTW, there's an Atari PC simulator: <https://stella-emu.github.io/> I spent a lot of time playing them...


C5 (10 pts): Complete the code for the CNN with 3 conv layers and 3 fc layers in class SimpleCNN in file Model.py And verify the output shape with the cell below.

```
In [12]: ## Test code
         from Model import SimpleCNN
         import torch
         net = SimpleCNN()
         x = torch.randn(2, 4, 84, 84)
         y = net(x)
         assert y.shape == (2, 4), "ERROR: network output has the wrong shape!"
         print ("CNN output shape test passed!")
```

CNN output shape test passed!


P3 (10 pts): Run the following cell to generate a DQN learning curve. The *maximum* average episodic reward on this curve should be larger than 200 for full credit. (It's ok if the final reward is not as high.) The typical value is around 300. You get 70% credit if $100 \leq \text{average episodic reward} < 200$, 50% credit if $50 \leq \text{average episodic reward} < 100$.

```
In [13]: plot_curve('log_breakout_dqn.txt', 'Breakout DQN')
```



P4 (10 pts): Run the following cell to generate an A2C learning curve. The *maximum* average episodic reward on this curve should be larger than 150 for full credit. (It's ok if the final reward is not as high.) The typical value is around 250. You get 70% credit if $50 \leq \text{average episodic reward} < 150$, and 50% credit if $20 \leq \text{average episodic reward} < 50$.

```
In [14]: plot_curve('log_breakout_a2c.txt', 'Breakout A2C')
```




P5 (10 pts): Collect and visualize some game frames by running the script `Draw.py` on BlueWaters.

- (1) module load python/2.0.0 and run `Draw.py` on BlueWaters (it's ok to run this locally, no need to start a job).
- (2) Download the result `breakout_imgs` folder from BlueWaters to the folder containing this Jupyter notebook, and run the following cell. You should see some animation of the game.

```
In [15]: import os
         imgs = sorted(os.listdir('breakout_imgs'))
         #imgs = [plt.imread('breakout_imgs/' + img) for img in imgs]

         %matplotlib inline
         import matplotlib.pyplot as plt
         from IPython import display
         pimg = None
         for img in imgs:
             img = plt.imread('breakout_imgs/' + img)
             if pimg:
                 pimg.set_data(img)
             else:
                 pimg = plt.imshow(img)
             display.display(plt.gcf())
             display.clear_output(wait=True)
```

1.3 Part III: Questions (10 pts)

These are open-ended questions. The purpose is to encourage you to think (a bit) more deeply about these algorithms. You get full points as long as you write a few sentences that make sense and show some thinking.

Q1 (2 pts): Why would people want to do function approximation rather than using tabular algorithm (on discretized S,A spaces if necessary)? Bringing function approximation has caused numerous problems theoretically (e.g. not guaranteed to converge), so it seems not worth it... Your answer: Deep Q learning as function approximator deals efficiently with curse of dimensionality. Using CNN as components of RL, we can directly learn from raw, high-dimensional visual inputs.

Q2 (2 pts): Q-Learning seems good... it's theoretically sound (at least seems to be), the performance is also good. Why would many people actually prefer policy gradient type algorithms in some practical problems? Your answer: Sometimes, Q function is too complex to be learned. Policy gradient would be still capable since it directly operates in the policy space. Also, Policy gradient usually converges faster and it may also learn the stochastic policies. Moreover, since the policy network is designed to model probability distribution, it is easy to apply to model continuous action space.

Q3 (2 pts): Does the policy gradient algorithm (A2C) we implemented here extend to continuous action space? How would you do that? Hint: What is a reasonable distribution assumption for policy $\pi_\theta(a|s)$ if a lives in continuous space? Your answer: Our implementation does not extend to continuous action space. To do that, instead of let the network output the parameters for a categorical distribution, we may output the parameters for a Gaussian distribution with mean and std.

Q4 (2 pts): The policy gradient algorithm (A2C) we implemented uses on-policy data. Can you think of a way to extend it to utilize off-policy data? Hint: Importance sampling, needs some approximation though Your answer: Utilizing off-policy data means that we can use any actions to improve your value/action-value functions instead of only using the actions generated by the policy. However, we would get a lot of samples that are not part of distribution that we are interested in, which causes high variances. To filter out those unrelated samples, we can use importance sampling with approximation to determine how important the samples generated are to samples that the target policy may have made.

Q5 (2 pts): How to compare different RL algorithms? When can I say one algorithm is better than the other? Hint: This question is quite open. Think about speed, complexity, tasks, etc. Your answer: - Considering sample efficiency from less to more efficient: on-policy policy gradient algorithms -> actor-critic style methods -> off-policy Q-function learning -> model-based.

- Value function fitting(Q-learning, DQN) may minimize error of fit at best and doesn't optimize anything at worst. The complexity could be very high. Q-learning uses fixed point iteration that may not converge.

- Model-based RL guarantees to converge and minimizes error of fit but does not guarantee that better model is better policy. It may not optimize for expected reward.

- Policy gradient is the only one that actually performs gradient descent on the true objective. But it is also often the least efficient.